

***L<sup>A</sup>T<sub>E</sub>X* command declarations here.**

In [1]:

```

from __future__ import division

# scientific
%matplotlib inline
from matplotlib import pyplot as plt;
import numpy as np;

# ipython
import IPython;

# python
import os;

#####

# image processing
import PIL;

# trim and scale images
def trim(im, percent=100):
    print("trim:", percent);
    bg = PIL.Image.new(im.mode, im.size, im.getpixel((0,0)))
    diff = PIL.ImageChops.difference(im, bg)
    diff = PIL.ImageChops.add(diff, diff, 2.0, -100)
    bbox = diff.getbbox()
    if bbox:
        x = im.crop(bbox)
        return x.resize(((x.size[0]*percent)//100, (x.size[1]*percent)//100), PIL.Image.ANTIALIAS);

#####

# daft (rendering PGMs)
import daft;

# set to FALSE to load PGMs from static images
RENDER_PGMS = False;

# decorator for pgm rendering
def pgm_render(pgm_func):
    def render_func(path, percent=100, render=None, *args, **kwargs):
        print("render_func:", percent);
        # render
        render = render if (render is not None) else RENDER_PGMS;

        if render:
            print("rendering");
            # render
            pgm = pgm_func(*args, **kwargs);
            pgm.render();

```

```

pgm.figure.savefig(path, dpi=300);

# trim
img = trim(PIL.Image.open(path), percent);
img.save(path, 'PNG');
else:
    print("not rendering");

# error
if not os.path.isfile(path):
    raise "Error: Graphical model image %s not found.
You may need to set RENDER_PGMS=True.";

# display
return IPython.display.Image(filename=path);#trim(PIL.Im
age.open(path), percent);

return render_func;

#####

```

## EECS 545: Machine Learning

### Lecture 14: Exponential Families & Bayesian Networks

- Instructor: **Jacob Abernethy**
- Date: March 9, 2016

*Lecture Exposition Credit:* Benjamin Bray & Valliappa Chockalingam

### References

- **[MLAPP]** Murphy, Kevin. *Machine Learning: A Probabilistic Perspective* (<https://mitpress.mit.edu/books/machine-learning-0>). 2012.
- **[Koller & Friedman 2009]** Koller, Daphne and Nir Friedman. *Probabilistic Graphical Models* (<https://mitpress.mit.edu/books/probabilistic-graphical-models>). 2009.
- **[Hero 2008]** Hero, Alfred O.. *Statistical Methods for Signal Processing* ([http://web.eecs.umich.edu/~hero/Preprints/main\\_564\\_08\\_new.pdf](http://web.eecs.umich.edu/~hero/Preprints/main_564_08_new.pdf)). 2008.
- **[Blei 2011]** Blei, David. *Notes on Exponential Families* (<https://www.cs.princeton.edu/courses/archive/fall11/cos597C/lectures/exponential-families.pdf>). 2011.
- **[Jordan 2010]** Jordan, Michael I.. *The Exponential Family: Basics* (<http://www.cs.berkeley.edu/~jordan/courses/260-spring10/other-readings/chapter8.pdf>). 2008.

# Outline

- Exponential Families
  - Sufficient Statistics & Pitman-Koopman-Darmois Theorem
  - Mean and natural parameters
  - Maximum Likelihood estimation
- Probabilistic Graphical Models
  - Directed Models (Bayesian Networks)
  - Conditional Independence & Factorization
  - Examples

## Exponential Families

Uses material from [MLAPP] §9.2 and [Hero 2008] §3.5, §4.4.2

### Exponential Family: Introduction

We have seen many distributions.

- Bernoulli
- Gaussian
- Exponential
- Gamma

Many of these belong to a more general class called the **exponential family**.

### Exponential Family: Introduction

Why do we care?

- only family of distributions with finite-dimensional **sufficient statistics**
- only family of distributions for which **conjugate priors** exist
- makes the least set of assumptions subject to some user-chosen constraints (**Maximum Entropy**)
- core of generalized linear models and **variational inference**

## Sufficient Statistics

**Recall:** A **statistic**  $T(\mathcal{D})$  is a function of the observed data  $\mathcal{D}$ .

- Mean,  $T(x_1, \dots, x_n) = \frac{1}{n} \sum_{k=1}^n x_k$
- Variance, maximum, mode, etc.

## Sufficient Statistics: Definition

Suppose we have a model  $P$  with parameters  $\theta$ . Then,

A statistic  $T(\mathcal{D})$  is **sufficient** for  $\theta$  if no other statistic calculated from the same sample provides any additional information about the parameter.

That is, if  $T(\mathcal{D}_1) = T(\mathcal{D}_2)$ , our estimate of  $\theta$  given  $\mathcal{D}_1$  or  $\mathcal{D}_2$  will be the same.

- Mathematically,  $P(\theta | T(\mathcal{D}), \mathcal{D}) = P(\theta | T(\mathcal{D}))$  independently of  $\mathcal{D}$

## Sufficient Statistics: Example

Suppose  $X \sim \mathcal{N}(\mu, \sigma^2)$  and we observe  $\mathcal{D} = (x_1, \dots, x_n)$ . Let

- $\hat{\mu}$  be the sample mean
- $\hat{\sigma}^2$  be the sample variance

Then  $T(\mathcal{D}) = (\hat{\mu}, \hat{\sigma}^2)$  is sufficient for  $\theta = (\mu, \sigma^2)$ .

- Two samples  $\mathcal{D}_1$  and  $\mathcal{D}_2$  with the same mean and variance give the same estimate of  $\theta$

(we are sweeping some details under the rug)

## Exponential Family: Definition

$p(x|\theta)$  has **exponential family form** if:

$$\begin{aligned} p(x|\theta) &= \frac{1}{Z(\theta)} h(x) \exp[\eta(\theta)^T \phi(x)] \\ &= h(x) \exp[\eta(\theta)^T \phi(x) - A(\theta)] \end{aligned}$$

- $Z(\theta)$  is the **partition function** for normalization
- $A(\theta) = \log Z(\theta)$  is the **log partition function**
- $\phi(x) \in \mathbb{R}^d$  is a vector of **sufficient statistics**
- $\eta(\theta)$  maps  $\theta$  to a set of **natural parameters**
- $h(x)$  is a scaling constant, usually  $h(x) = 1$

## Example: Bernoulli

The Bernoulli distribution can be written as

$$\begin{aligned} \text{Ber}(x|\mu) &= \mu^x (1 - \mu)^{1-x} \\ &= \exp[x \log \mu + (1 - x) \log(1 - \mu)] \\ &= \exp[\eta(\mu)^T \phi(x)] \end{aligned}$$

where  $\eta(\mu) = (\log \mu, \log(1 - \mu))$  and  $\phi(x) = (x, 1 - x)$

- There is a linear dependence between features  $\phi(x)$
- This representation is **overcomplete**
- $\eta$  is not uniquely determined

## Example: Bernoulli

Instead, we can find a **minimal** parameterization:

$$\text{Ber}(x|\mu) = (1 - \mu) \exp\left[x \log \frac{\mu}{1 - \mu}\right]$$

This gives **natural parameters**  $\eta = \log \frac{\mu}{1 - \mu}$ .

- Now,  $\eta$  is unique

## Other Examples

Exponential Family Distributions:

- Multivariate normal
- Exponential
- Dirichlet

Non-examples:

- Student t-distribution can't be written in exponential form
- Uniform distribution support depends on the parameters  $\theta$

## Log-Partition Function

Derivatives of the **log-partition function**  $A(\theta)$  yield **cumulants** of the sufficient statistics (*Exercise!*)

- $\nabla_{\theta} A(\theta) = E[\phi(x)]$
- $\nabla_{\theta}^2 A(\theta) = \text{Cov}[\phi(x)]$

This guarantees that  $A(\theta)$  is convex!

- Its Hessian is the covariance matrix of  $\mathbf{X}$ , which is positive-definite.
- Later, this will guarantee a unique global maximum of the likelihood!

**Proof of Convexity: First Derivative**

$$\begin{aligned}
 \frac{dA}{d\theta} &= \frac{d}{d\theta} \left[ \log \int \exp(\theta \phi(x)) h(x) dx \right] \\
 &= \frac{\frac{d}{d\theta} \int \exp(\theta \phi(x)) h(x) dx}{\int \exp(\theta \phi(x)) h(x) dx} \\
 &= \frac{\int \phi(x) \exp(\theta \phi(x)) h(x) dx}{\exp(A(\theta))} \\
 &= \int \phi(x) \exp[\theta \phi(x) - A(\theta)] h(x) dx \\
 &= \int \phi(x) p(x) dx \\
 &= E[\phi(x)]
 \end{aligned}$$



**Proof of Convexity: Second Derivative**

$$\begin{aligned}
\frac{d^2 A}{d\theta^2} &= \int \phi(x) \exp[\theta\phi(x) - A(\theta)] h(x) (\phi(x) - A'(\theta)) dx \\
&= \int \phi(x) p(x) (\phi(x) - A'(\theta)) dx \\
&= \int \phi^2(x) p(x) dx - A'(\theta) \int \phi(x) p(x) dx \\
&= E[\phi^2(x)] - E[\phi(x)]^2 \quad (\because A'(\theta) = E[\phi(x)]) \\
&= \text{Var}[\phi(x)]
\end{aligned}$$

**Proof of Convexity: Second Derivative**

For multi-variate case, we have

$$\frac{\partial^2 A}{\partial \theta_i \partial \theta_j} = E[\phi_i(x) \phi_j(x)] - E[\phi_i(x)] E[\phi_j(x)]$$

and hence,

$$\nabla^2 A(\theta) = \text{Cov}[\phi(x)]$$

Since covariance is positive definite, we have  $A(\theta)$  convex as required.

**Exponential Family: Likelihood**

For data  $\mathcal{D} = (x_1, \dots, x_N)$ , the likelihood is

$$p(\mathcal{D}|\theta) = \left[ \prod_{k=1}^N h(x_k) \right] Z(\theta)^{-N} \exp \left[ \eta(\theta)^T \left( \sum_{k=1}^N \phi(x_k) \right) \right]$$

The sufficient statistics are now  $\phi(\mathcal{D}) = \sum_{k=1}^N \phi(x_k)$ .

- **Bernoulli:**  $\phi = \#Heads$
- **Normal:**  $\phi = [\sum_k x_k, \sum_k x_k^2]$

## Pitman-Koopman-Darmois Theorem

Among families of distributions  $P(\mathbf{x}|\theta)$  whose support does not vary with the parameter  $\theta$ , only in exponential families is there a sufficient statistic  $T(\mathbf{x}_1, \dots, \mathbf{x}_N)$  whose dimension remains bounded as the sample size  $N$  increases.

## Exponential Family: MLE

For natural parameters  $\theta$  and data  $\mathcal{D} = (\mathbf{x}_1, \dots, \mathbf{x}_N)$ ,

$$\log p(\mathcal{D}|\theta) = \theta^T \phi(\mathcal{D}) - NA(\theta)$$

Since  $-A(\theta)$  is concave and  $\theta^T \phi(\mathcal{D})$  linear,

- the log-likelihood is concave
- there is a unique global maximum!

## Exponential Family: MLE

To find the maximum, recall  $\nabla_{\theta} A(\theta) = E_{\theta}[\phi(\mathbf{x})]$ , so 
$$\nabla_{\theta} \log p(\mathcal{D} | \theta) = \nabla_{\theta} (\theta^T \phi(\mathcal{D}) - NA(\theta)) = \phi(\mathcal{D}) - NE_{\theta}[\phi(\mathbf{x})] = 0$$
 Which gives

$$E_{\theta}[\phi(\mathbf{x})] = \frac{\phi(\mathcal{D})}{N} = \frac{1}{N} \sum_{k=1}^N \phi(\mathbf{x}_k)$$

At the MLE  $\hat{\theta}_{MLE}$ , the empirical average of sufficient statistics equals their expected value.

- this is called **moment matching**

## Exponential Family: MLE

As an example, consider the Bernoulli distribution

- Sufficient statistic  $N, \phi(\mathcal{D}) = \#Heads$

$$\hat{\mu}_{MLE} = \frac{\#Heads}{N}$$

## Bayes for Exponential Family

Exact Bayesian analysis is considerably simplified if the prior is **conjugate** to the likelihood.

- Simply, this means that prior  $p(\theta)$  has the same form as the posterior  $p(\theta|\mathcal{D})$ .

This requires likelihood to have finite sufficient statistics

- Exponential family to the rescue!

**Note:** We will release some notes on conjugate priors + exponential families. It's hard to learn from slides and needs a bit more description.

## Likelihood for exponential family

Likelihood:

$$p(\mathcal{D}|\theta) \propto g(\theta)^N \exp[\eta(\theta)^T s_N]$$

$$s_N = \sum_{i=1}^N \phi(x_i)$$

In terms of canonical parameters:

$$p(\mathcal{D}|\eta) \propto \exp[N\eta^T \bar{s} - NA(\eta)]$$

$$\bar{s} = \frac{1}{N} s_N$$

## Conjugate prior for exponential family

- The prior and posterior for an exponential family involve two parameters,  $\tau$  and  $\nu$ , initially set to  $\tau_0, \nu_0$

$$p(\theta|\nu_0, \tau_0) \propto g(\theta)^{\nu_0} \exp[\eta(\theta)^T \tau_0]$$

- Denote  $\tau_0 = \nu_0 \bar{\tau}_0$  to separate out the size of the **prior pseudo-data**,  $\nu_0$ , from the mean of the sufficient statistics on this pseudo-data,  $\tau_0$ . Hence,

$$p(\theta|\nu_0, \bar{\tau}_0) \propto \exp[\nu_0 \eta^T \bar{\tau}_0 - \nu_0 A(\eta)]$$

- Think of  $\tau_0$  as a "guess" of the future sufficient statistics, and  $\nu_0$  as the strength of this guess

## Prior: Example

$$\begin{aligned} p(\theta|\nu_0, \tau_0) &\propto (1 - \theta)^{\nu_0} \exp[\tau_0 \log(\frac{\theta}{1 - \theta})] \\ &= \theta^{\tau_0} (1 - \theta)^{\nu_0 - \tau_0} \end{aligned}$$

Define  $\alpha = \tau_0 + 1$  and  $\beta = \nu_0 - \tau_0 + 1$  to see that this is a **beta distribution**.

## Posterior

Posterior:

$$p(\theta|\mathcal{D}) = p(\theta|\nu_N, \tau_N) = p(\theta|\nu_0 + N, \tau_0 + s_N)$$

Note that we obtain **hyper-parameters** by adding. Hence,

$$\begin{aligned} p(\eta|\mathcal{D}) &\propto \exp[\eta^T (\nu_0 \bar{\tau}_0 + N \bar{s}) - (\nu_0 + N) A(\eta)] \\ &= p(\eta|\nu_0 + N, \frac{\nu_0 \bar{\tau}_0 + N \bar{s}}{\nu_0 + N}) \end{aligned}$$

where  $\bar{s} = \frac{1}{N} \sum_{i=1}^N \phi(x_i)$ .

- *posterior hyper-parameters are a convex combination of the prior mean hyper-parameters and the average of the sufficient statistics.*

## Break time!



# Probabilistic Graphical Models

Uses material from [MLAPP] §10.1, 10.2 and [Koller & Friedman 2009].

"I basically know of two principles for treating complicated systems in simple ways: the first is the principle of modularity and the second is the principle of abstraction. I am an apologist for computational probability in machine learning because I believe that probability theory implements these two principles in deep and intriguing ways — namely through factorization and through averaging. Exploiting these two mechanisms as fully as possible seems to me to be the way forward in machine learning" — Michael Jordan (qtd. in MLAPP)

## Graphical Models: Motivation

Suppose we observe multiple correlated variables  $\mathbf{x} = (x_1, \dots, x_n)$ .

- Words in a document
- Pixels in an image

How can we compactly represent the **joint distribution**  $p(\mathbf{x}|\theta)$ ?

- How can we tractably *infer* one set of variables given another?
- How can we efficiently *learn* the parameters?

## Joint Probability Tables

One (bad) choice is to write down a **Joint Probability Table**.

- For  $n$  binary variables, we must specify  $2^n - 1$  probabilities!
- Expensive to store and manipulate
- Impossible to learn so many parameters
- Very hard to interpret!

Can we be more concise?

## Motivating Example: Coin Flips

What is the joint distribution of three independent coin flips?

- Explicitly specifying the JPT requires  $2^3 - 1 = 7$  parameters.

Assuming independence,  $P(X_1, X_2, X_3) = P(X_1)P(X_2)P(X_3)$

- Each marginal  $P(X_k)$  only requires one parameter, the bias
- This gives a total of 3 parameters, compared to 8.

Exploiting the **independence structure** of a joint distribution leads to more concise representations.

## Motivating Example: Naive Bayes

In Naive Bayes, we assumed the features  $X_1, \dots, X_N$  were independent given the class label  $C$ :

$$P(x_1, \dots, x_N, c) = P(c) \prod_{k=1}^N P(x_k | c)$$

This greatly simplified the learning procedure:

- Allowed us to look at each feature individually
- Only need to learn  $O(CN)$  probabilities, for  $C$  classes and  $N$  features

## Conditional Independence

The key to efficiently representing large joint distributions is to make **conditional independence** assumptions of the form

$$X \perp Y \mid Z \iff p(X, Y | Z) = p(X | Z)p(Y | Z)$$

Once  $z$  is known, information about  $x$  does not tell us any information about  $y$  and vice versa.

An effective way to represent these assumptions is with a **graph**.

## Bayesian Networks: Definition

A **Bayesian Network**  $\mathcal{G}$  is a directed acyclic graph whose nodes represent random variables  $X_1, \dots, X_n$ .

- Let  $\text{Parents}_{\mathcal{G}}(X_k)$  denote the parents of  $X_k$  in  $\mathcal{G}$
- Let  $\text{NonDesc}_{\mathcal{G}}(X_k)$  denote the variables in  $\mathcal{G}$  who are not descendants of  $X_k$ .

Examples will come shortly...

## Bayesian Networks: Local Independencies

Every Bayesian Network  $\mathcal{G}$  encodes a set  $\mathcal{I}_{\ell}(\mathcal{G})$  of **local independence assumptions**:

For each variable  $X_k$ , we have  $(X_k \perp \text{NonDesc}_{\mathcal{G}}(X_k) \mid \text{Parents}_{\mathcal{G}}(X_k))$

Every node  $X_k$  is conditionally independent of its nondescendants given its parents.

## Example: Naive Bayes

The graphical model for Naive Bayes is shown below:

- $\text{Parents}_{\mathcal{G}}(X_k) = \{C\}$ ,  $\text{NonDesc}_{\mathcal{G}}(X_k) = \{X_j\}_{j \neq k}$
- Therefore  $X_j \perp X_k \mid C$  for any  $j \neq k$

```
In [2]: @pgm_render
def pgm_naive_bayes():
    pgm = daft.PGM([4,3], origin=[-2,0], node_unit=0.8, grid_unit=2.0);
    # nodes
    pgm.add_node(daft.Node("c", r"$C$", -0.25, 2));
    pgm.add_node(daft.Node("x1", r"$X_1$", -1, 1));
    pgm.add_node(daft.Node("x2", r"$X_2$", -0.5, 1));
    pgm.add_node(daft.Node("dots", r"$\cdots$", 0, 1, plot_params={ 'ec' : 'none' }));
    pgm.add_node(daft.Node("xN", r"$X_N$", 0.5, 1));

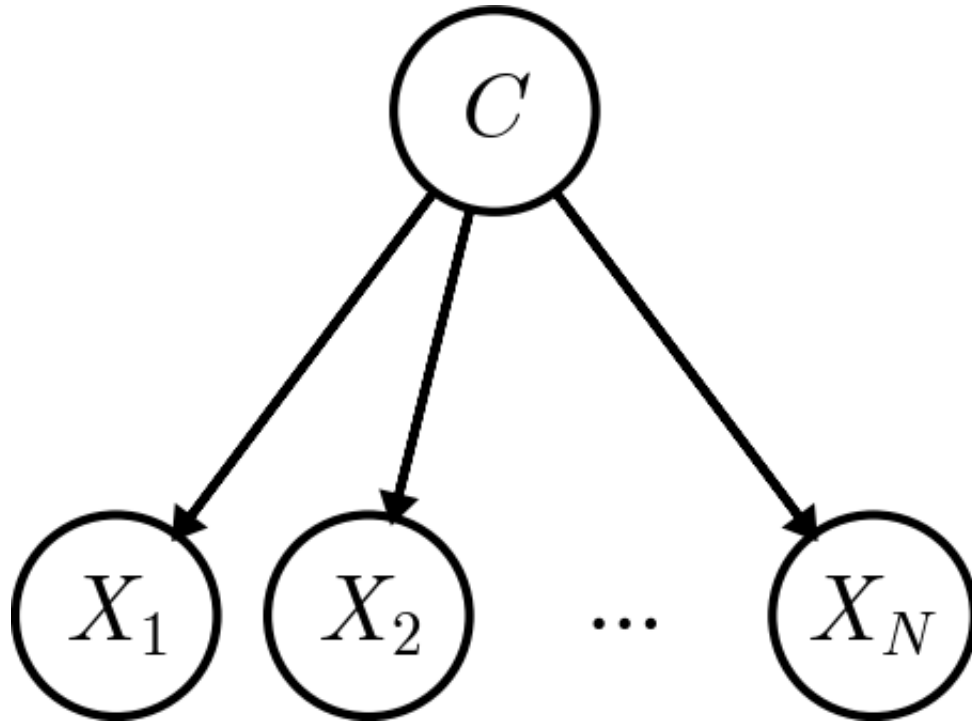
    # edges
    pgm.add_edge("c", "x1", head_length=0.08);
    pgm.add_edge("c", "x2", head_length=0.08);
    pgm.add_edge("c", "xN", head_length=0.08);

    return pgm;
```



```
In [3]: %%capture
pgm_naive_bayes("images/naive-bayes.png");
```

Out[3]:



## Subtle Point: Graphs & Distributions

A Bayesian network  $\mathcal{G}$  over variables  $X_1, \dots, X_N$  encodes a set of **conditional independencies**.

- Shows independence structure, nothing more.
- Does **not** tell us how to assign probabilities to a configuration  $(x_1, \dots, x_N)$  of the variables.

There are **many** distributions  $P$  satisfying the independencies in  $\mathcal{G}$ .

- Many joint distributions share a common structure, which we exploit in algorithms.
- The distribution  $P$  may satisfy other independencies **not** encoded in  $\mathcal{G}$ .

## Subtle Point: Graphs & Distributions

If  $P$  satisfies the independence assertions made by  $\mathcal{G}$ , we say that

- $\mathcal{G}$  is an **I-Map** for  $P$
- or that  $P$  **satisfies**  $\mathcal{G}$ .

Any distribution satisfying  $\mathcal{G}$  shares common structure.

- We will exploit this structure in our algorithms
- This is what makes graphical models so **powerful**!

## Review: Chain Rule for Probability

We can factorize any joint distribution via the **Chain Rule for Probability**:

$$\begin{aligned} P(X_1, \dots, X_N) &= P(X_1)P(X_2, \dots, X_N|X_1) \\ &= P(X_1)P(X_2|X_1)P(X_3, \dots, X_N|X_1, X_2) \\ &= \prod_{k=1}^N P(X_k|X_1, \dots, X_{k-1}) \end{aligned}$$

Here, the ordering of variables is arbitrary. This works for any permutation.

## Bayesian Networks: Topological Ordering

Every network  $\mathcal{G}$  induces a **topological (partial) ordering** on its nodes:

Parents assigned a lower index than their children

```
In [4]: @pgm_render
def pgm_topological_order():
    pgm = daft.PGM([4, 4], origin=[-4, 0])

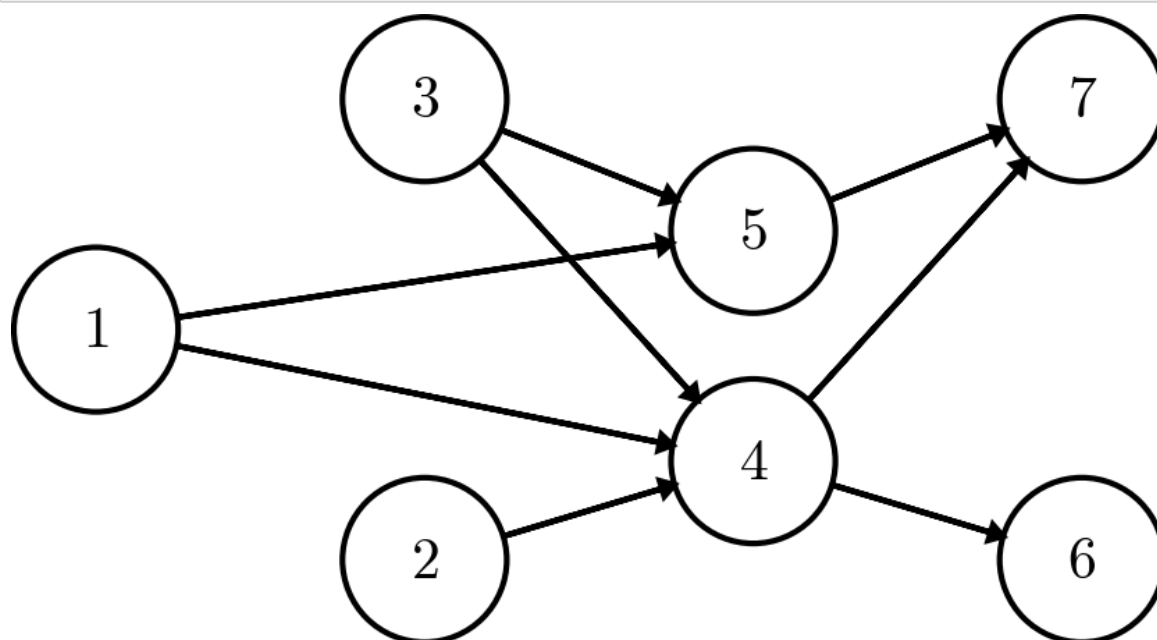
    # Nodes
    pgm.add_node(daft.Node("x1", r"$1$", -3.5, 2))
    pgm.add_node(daft.Node("x2", r"$2$", -2.5, 1.3))
    pgm.add_node(daft.Node("x3", r"$3$", -2.5, 2.7))
    pgm.add_node(daft.Node("x4", r"$4$", -1.5, 1.6))
    pgm.add_node(daft.Node("x5", r"$5$", -1.5, 2.3))
    pgm.add_node(daft.Node("x6", r"$6$", -0.5, 1.3))
    pgm.add_node(daft.Node("x7", r"$7$", -0.5, 2.7))

    # Add in the edges.
    pgm.add_edge("x1", "x4", head_length=0.08)
    pgm.add_edge("x1", "x5", head_length=0.08)
    pgm.add_edge("x2", "x4", head_length=0.08)
    pgm.add_edge("x3", "x4", head_length=0.08)
    pgm.add_edge("x3", "x5", head_length=0.08)
    pgm.add_edge("x4", "x6", head_length=0.08)
    pgm.add_edge("x4", "x7", head_length=0.08)
    pgm.add_edge("x5", "x7", head_length=0.08)

    return pgm;
```

```
In [5]: %%capture
pgm_topological_order("images/topological-order.png")
```

Out[5]:



## Factorization Theorem: Statement

**Theorem:** (Koller & Friedman 3.1) If  $\mathcal{G}$  is an I-map for  $P$ , then  $P$  factorizes as follows:

$$P(X_1, \dots, X_N) = \prod_{k=1}^N P(X_k \mid \text{Parents}_{\mathcal{G}}(X_k))$$

Let's prove it together!

## Factorization Theorem: Proof

First, apply the chain rule to any topological ordering:

$$P(X_1, \dots, X_N) = \prod_{k=1}^N P(X_k \mid X_1, \dots, X_{k-1})$$

Consider one of the factors  $P(X_k \mid X_1, \dots, X_{k-1})$ .

## Factorization Theorem: Proof

Since our variables  $X_1, \dots, X_N$  are in topological order,

- $\text{Parents}_{\mathcal{G}}(X_k) \subseteq \{X_1, \dots, X_{k-1}\}$
- None of  $X_k$ 's descendants can possibly lie in  $\{X_1, \dots, X_{k-1}\}$

Therefore,  $\{X_1, \dots, X_{k-1}\} = \text{Parents}_{\mathcal{G}}(X_k) \cup \mathcal{Z}$

- for some  $\mathcal{Z} \subseteq \text{NonDesc}_{\mathcal{G}}(X_k)$ .

## Factorization Theorem: Proof

Recall the following property of conditional independence:

$$(X \perp Y, W \mid Z) \implies (X \perp Y \mid Z)$$

Since  $\mathcal{G}$  is an I-map for  $P$  and  $\mathcal{Z} \subseteq \text{NonDesc}_{\mathcal{G}}(X_k)$ , we have

$$\begin{aligned} & (X_k \perp \text{NonDesc}_{\mathcal{G}}(X_k) \mid \text{Parents}_{\mathcal{G}}(X_k)) \\ \implies & (X_k \perp \mathcal{Z} \mid \text{Parents}_{\mathcal{G}}(X_k)) \end{aligned}$$

## Factorization Theorem: Proof

We have just shown  $(X_k \perp \mathcal{Z} \mid \text{Parents}_{\mathcal{G}}(X_k))$ , therefore

$$P(X_k \mid X_1, \dots, X_{k-1}) = P(X_k \mid \text{Parents}_{\mathcal{G}}(X_k))$$

- Recall  $\{X_1, \dots, X_N\} = \text{Parents}_{\mathcal{G}}(X_k) \cup \mathcal{Z}$ .

**Remember:**  $X_k$  is conditionally independent of its nondescendants given its parents!

## Factorization Theorem: End of Proof

Applying this to every factor, we see that

$$\begin{aligned} P(X_1, \dots, X_N) &= \prod_{k=1}^N P(X_k \mid X_1, \dots, X_{k-1}) \\ &= \prod_{k=1}^N P(X_k \mid \text{Parents}_{\mathcal{G}}(X_k)) \end{aligned}$$

## Factorization Theorem: Consequences

We just proved that for any  $P$  satisfying  $\mathcal{G}$ ,

$$P(X_1, \dots, X_N) = \prod_{k=1}^N P(X_k \mid \text{Parents}_{\mathcal{G}}(X_k))$$

It suffices to store **conditional probability tables**  $P(X_k \mid \text{Parents}_{\mathcal{G}}(X_k))$ !

- Requires  $O(N2^k)$  features if each node has  $\leq k$  parents
- Substantially more compact than **JPTs** for  $N$  large,  $\mathcal{G}$  sparse
- We can also specify that a CPD is Gaussian, Dirichlet, etc.

## Example: Fully Connected Graph

A **fully connected graph** makes no independence assumptions.

$$P(A, B, C) = P(A)P(B|A)P(C|A, B)$$

```
In [6]: @pgm_render
def pgm_fully_connected_a():
    pgm = daft.PGM([4, 4], origin=[0, 0])

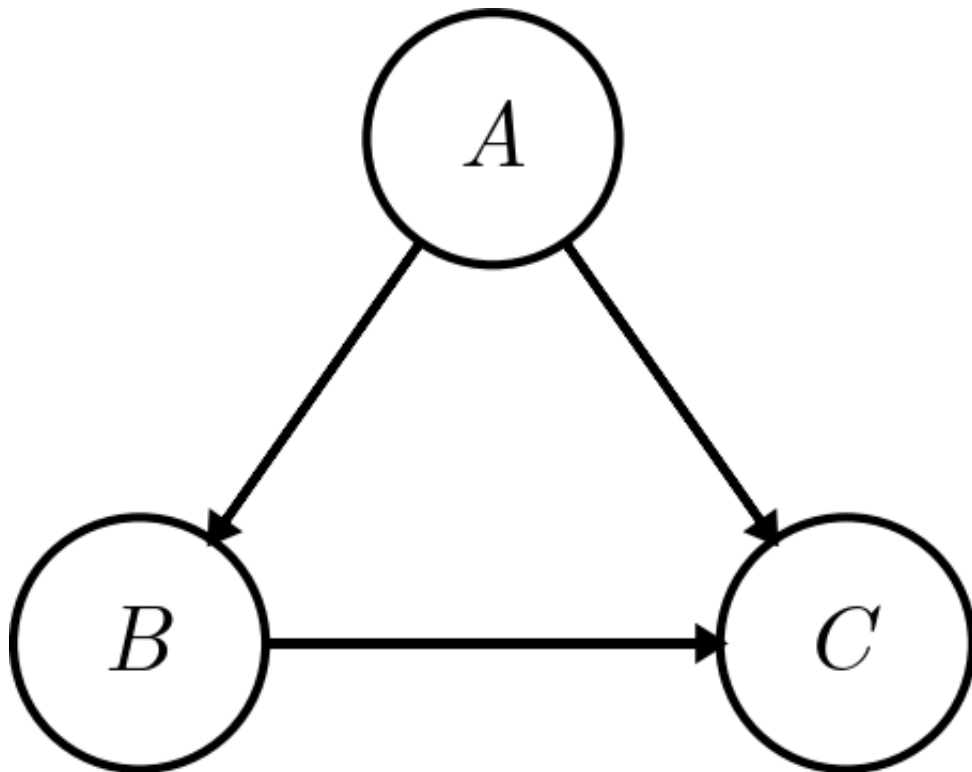
    # nodes
    pgm.add_node(daft.Node("a", r"$A$", 2, 3.5))
    pgm.add_node(daft.Node("b", r"$B$", 1.3, 2.5))
    pgm.add_node(daft.Node("c", r"$C$", 2.7, 2.5))

    # add in the edges
    pgm.add_edge("a", "b", head_length=0.08)
    pgm.add_edge("a", "c", head_length=0.08)
    pgm.add_edge("b", "c", head_length=0.08)

    return pgm;
```

```
In [7]: %%capture
pgm_fully_connected_a("images/fully-connected-a.png")
```

Out[7]:



## Example: Fully Connected Graph

There are many possible fully connected graphs:

$$\begin{aligned} P(A, B, C) &= P(A)P(B|A)P(C|A, B) \\ &= P(B)P(C|B)P(A|B, C) \end{aligned}$$

```
In [8]: @pgm_render
def pgm_fully_connected_b():
    pgm = daft.PGM([8, 4], origin=[0, 0])

    # nodes
    pgm.add_node(daft.Node("a1", r"$A$", 2, 3.5))
    pgm.add_node(daft.Node("b1", r"$B$", 1.5, 2.8))
    pgm.add_node(daft.Node("c1", r"$C$", 2.5, 2.8))

    # add in the edges
    pgm.add_edge("a1", "b1", head_length=0.08)
    pgm.add_edge("a1", "c1", head_length=0.08)
    pgm.add_edge("b1", "c1", head_length=0.08)

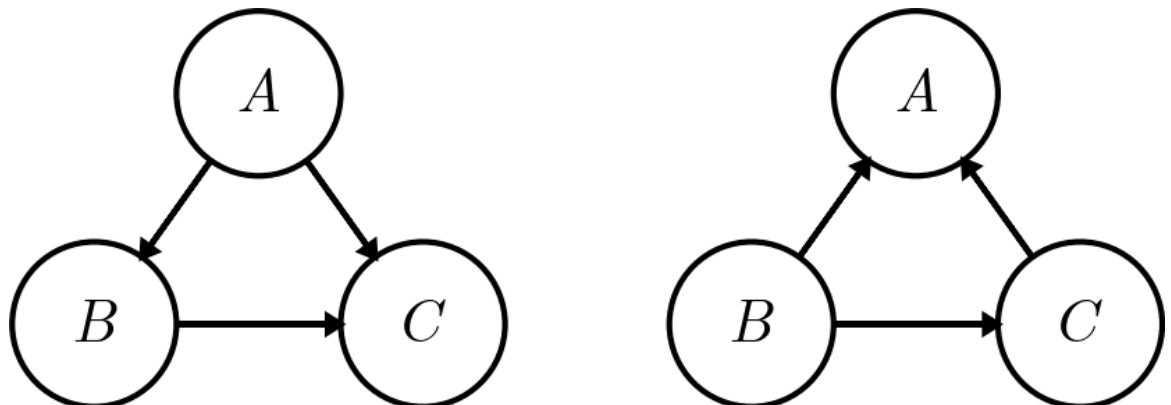
    # nodes
    pgm.add_node(daft.Node("a2", r"$A$", 4, 3.5))
    pgm.add_node(daft.Node("b2", r"$B$", 3.5, 2.8))
    pgm.add_node(daft.Node("c2", r"$C$", 4.5, 2.8))

    # add in the edges
    pgm.add_edge("b2", "c2", head_length=0.08)
    pgm.add_edge("b2", "a2", head_length=0.08)
    pgm.add_edge("c2", "a2", head_length=0.08)

    return pgm;
```

```
In [9]: %%capture
pgm_fully_connected_b("images/fully-connected-b.png")
```

Out[9]:



## Bayesian Networks & Causality

The fully-connected example brings up a crucial point:

Directed edges do **not** necessarily represent causality.

Bayesian networks encode **independence assumptions** only.

- This representation is not unique.

## Example: Markov Chain

State at time  $t$  depends only on state at time  $t - 1$ .

$$P(X_0, X_1, \dots, X_N) = P(X_0) \prod_{t=1}^N P(X_t | X_{t-1})$$

```
In [10]: @pgm_render
def pgm_markov_chain():
    pgm = daft.PGM([6, 6], origin=[0, 0])

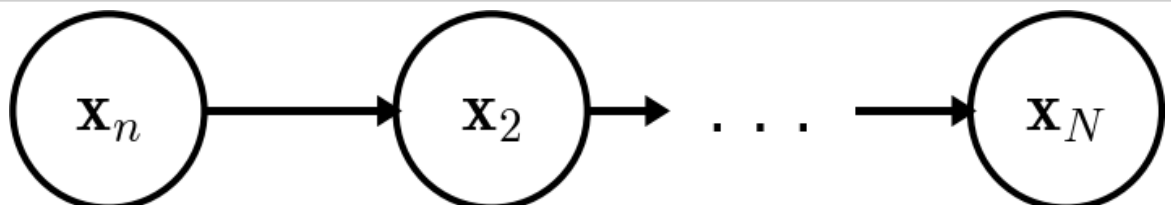
    # Nodes
    pgm.add_node(daft.Node("x1", r"$\mathbf{x}_n$", 2, 2.5))
    pgm.add_node(daft.Node("x2", r"$\mathbf{x}_2$", 3, 2.5))
    pgm.add_node(daft.Node("ellipsis", r" . . . ", 3.7, 2.5, offset=(0, 0), plot_params={"ec" : "none"}))
    pgm.add_node(daft.Node("ellipsis_end", r"", 3.7, 2.5, offset=(0, 0), plot_params={"ec" : "none"}))
    pgm.add_node(daft.Node("xN", r"$\mathbf{x}_N$", 4.5, 2.5))

    # Add in the edges.
    pgm.add_edge("x1", "x2", head_length=0.08)
    pgm.add_edge("x2", "ellipsis", head_length=0.08)
    pgm.add_edge("ellipsis_end", "xN", head_length=0.08)

    return pgm;
```

```
In [11]: %%capture
pgm_markov_chain("images/markov-chain.png")
```

Out[11]:





## Example: Hidden Markov Model

Noisy observations  $X_k$  generated from hidden Markov chain  $Y_k$ .

$$P(\mathbf{X}, \mathbf{Y}) = P(Y_1)P(X_1 | Y_1) \prod_{k=2}^N (P(Y_k | Y_{k-1})P(X_k | Y_k))$$

```
In [12]: @pgm_render
def pgm_hmm():
    pgm = daft.PGM([7, 7], origin=[0, 0])

    # Nodes
    pgm.add_node(daft.Node("Y1", r"$Y_1$", 1, 3.5))
    pgm.add_node(daft.Node("Y2", r"$Y_2$", 2, 3.5))
    pgm.add_node(daft.Node("Y3", r"$\dots$", 3, 3.5, plot_params
= {'ec': 'none'}))
    pgm.add_node(daft.Node("Y4", r"$Y_N$", 4, 3.5))

    pgm.add_node(daft.Node("x1", r"$X_1$", 1, 2.5, observed=True
e))
    pgm.add_node(daft.Node("x2", r"$X_2$", 2, 2.5, observed=True
e))
    pgm.add_node(daft.Node("x3", r"$\dots$", 3, 2.5, plot_params
= {'ec': 'none'}))
    pgm.add_node(daft.Node("x4", r"$X_N$", 4, 2.5, observed=True
e))

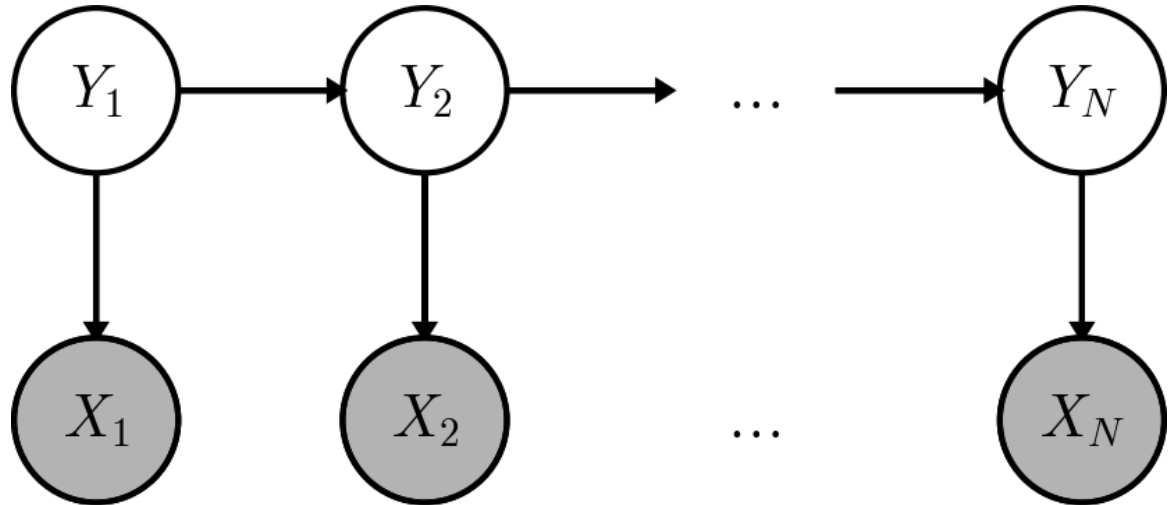
    # Add in the edges.
    pgm.add_edge("Y1", "Y2", head_length=0.08)
    pgm.add_edge("Y2", "Y3", head_length=0.08)
    pgm.add_edge("Y3", "Y4", head_length=0.08)

    pgm.add_edge("Y1", "x1", head_length=0.08)
    pgm.add_edge("Y2", "x2", head_length=0.08)
    pgm.add_edge("Y4", "x4", head_length=0.08)

    return pgm;
```

```
In [13]: %%capture  
pgm_hmm("images/hmm.png")
```

Out[13]:



### Example: Plate Notation

We can represent (conditionally) iid variables using **plate notation**.

```
In [14]: @pgm_render
def pgm_plate_example():
    pgm = daft.PGM([4,3], origin=[-2,0], node_unit=0.8, grid_uni
t=2.0);
    # nodes
    pgm.add_node(daft.Node("lambda", r"$\lambda$", -0.25, 2));

    pgm.add_node(daft.Node("t1", r"$\theta_1$", -1, 1.3));
    pgm.add_node(daft.Node("t2", r"$\theta_2$", -0.5, 1.3));
    pgm.add_node(daft.Node("dots1", r"$\cdots$", 0, 1.3, plot_p
arams={ 'ec' : 'none' }));
    pgm.add_node(daft.Node("tN", r"$\theta_N$", 0.5, 1.3));

    pgm.add_node(daft.Node("x1", r"$X_1$", -1, 0.6));
    pgm.add_node(daft.Node("x2", r"$X_2$", -0.5, 0.6));
    pgm.add_node(daft.Node("dots2", r"$\cdots$", 0, 0.6, plot_p
arams={ 'ec' : 'none' }));
    pgm.add_node(daft.Node("xN", r"$X_N$", 0.5, 0.6));


    pgm.add_node(daft.Node("LAMBDA", r"$\lambda$", 1.5, 2));
    pgm.add_node(daft.Node("THETA", r"$\theta_k$", 1.5,1.3));
    pgm.add_node(daft.Node("XX", r"$X_k$", 1.5,0.6));


    # edges
    pgm.add_edge("lambda", "t1", head_length=0.08);
    pgm.add_edge("lambda", "t2", head_length=0.08);
    pgm.add_edge("lambda", "tN", head_length=0.08);
    pgm.add_edge("t1", "x1", head_length=0.08);
    pgm.add_edge("t2", "x2", head_length=0.08);
    pgm.add_edge("tN", "xN", head_length=0.08);


    pgm.add_edge("LAMBDA", "THETA", head_length=0.08);
    pgm.add_edge("THETA", "XX", head_length=0.08);


    pgm.add_plate(daft.Plate([1.1,0.4,0.8,1.2], label=r"$\quad
\quad$; K$",
shift=-0.1))

    return pgm;
```

```
In [15]: %%capture  
pgm_plate_example("images/plate-example.png")
```

Out[15]:

